Machine Learning — Statistical Methods for Machine Learning

Consistency and nonparametric algorithms

Instructor: *Nicolò Cesa-Bianchi* version of March 30, 2023

Consistency is an asymptotical property certifying that the risk of the predictors generated by a learning algorithm converges to the Bayes risk in expectation as the size of the training set increases. Recall that $A(S_m)$ is the predictor generated by a learning algorithm A on a training set S_m of size m. A learning algorithm A is **consistent** with respect to a loss function ℓ if for any data distribution $\mathcal{D}$ it holds that

$$\lim_{m\to\infty} \mathbb{E}\Big[\ell_{\mathcal{D}}\big(A(S_m)\big)\Big] = \ell_{\mathcal{D}}(f^*)$$

where the expectation is with respect to the random draw of the training set S_m of size m from the distribution $\mathcal{D}$, and $\ell_{\mathcal{D}}(f^*)$ is the Bayes risk for $(\mathcal{D}, \ell)$. In some cases, we may define consistency with respect to a restricted class of distributions $\mathcal{D}$. For example, in binary classification we may restrict to all distributions $\mathcal{D}$ such that $\eta(\boldsymbol{x}) = \mathbb{P}(Y = 1 \mid \boldsymbol{X} = \boldsymbol{x})$ is a Lipschitz function on $\mathcal{X}$. Formally, there exists $0 < c < \infty$ such that

$$\big|\eta(\boldsymbol{x}) - \eta(\boldsymbol{x}')\big| \le c\,\big\|\boldsymbol{x} - \boldsymbol{x}'\big\| \qquad \text{for all } \boldsymbol{x}, \boldsymbol{x}' \in \mathcal{X} \ .$$

Technically, this conditions implies that η is Lipschitz in $\mathcal{X}$. This is a restriction on the set of all allowed η as $c < \infty$ implies continuity (but the opposite is not true).

Nonparametric algorithms. Given a learning algorithm A, let $\mathcal{H}_m$ be the set of predictors generated by A on training sets of size m: $h \in \mathcal{H}_m$ if and only if there exists a training set S_m of size m such that $A(S_m) = h$. We say that A is a nonparametric learning algorithm if A's approximation error vanishes as m grows to infinity. Formally,

$$\lim_{m\to\infty} \min_{h\in\mathcal{H}_m} \ell_{\mathcal{D}}(h^*) = \ell_{\mathcal{D}}(f^*) \ .$$

Two notable examples of nonparametric learning algorithms are k-NN and the greedy algorithm for decision tree classifiers (i.e., the algorithm that always chooses to split a leaf that maximizes the decrease in training error). Nonparametric algorithms are recognizable because:

- the size (memory footprint) of their predictors tends to grow with the training set size
- for any m and for all S_m, $\min_{h\in\mathcal{H}_m} \ell_S(s)$ is close to zero.

The standard k-NN algorithm is nonparametric but not known to be consistent for any fixed value of k. Indeed, one can only show that

$$\lim_{m\to\infty} \mathbb{E}\Big[\ell_{\mathcal{D}}\big(k\text{-NN}(S_m)\big)\Big] \le \ell_{\mathcal{D}}(f^*) + 2\sqrt{\frac{\ell_{\mathcal{D}}(f^*)}{k}} \tag{1}$$

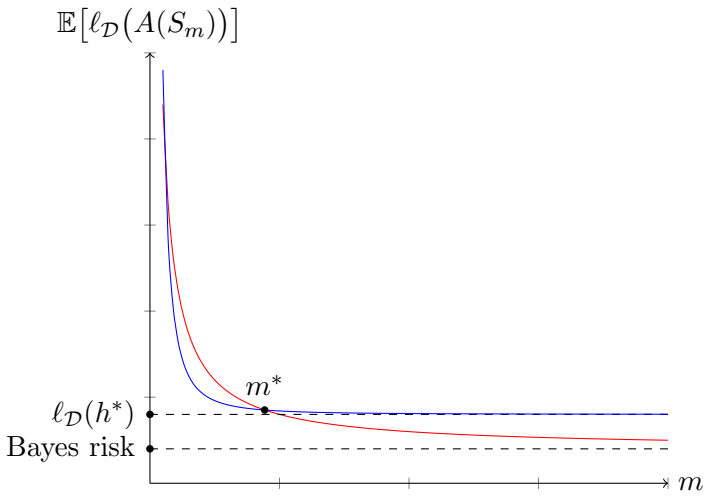

Figure 1: Typical behavior of expected risk $\mathbb{E}\big[\ell_{\mathcal{D}}\big(A(S_m)\big)\big]$ as a function of training set size for a consistent algorithm (red line) and for a nonconsistent algorithm (blue line). For small training set sizes $m < m^*$, the nonconsistent algorithm has a better performance. (Thanks to Edoardo Marangoni for drawing the picture.)

for any data distribution $\mathcal{D}$. However, if we let k be chosen as a function k_m of the training set size, then the algorithm becomes consistent provided $k_m \to \infty$ and $k_m = o(m)$.

Similarly, the greedy algorithm for building tree classifiers is consistent (for $\mathcal{X} \equiv \mathbb{R}^d$) whenever the two following conditions are fulfilled. Let $\ell(\boldsymbol{x})$ be the leaf to which $\boldsymbol{x} \in \mathbb{R}^d$ is routed in the current tree and let N_ℓ be the number of training examples routed to a leaf ℓ. Then, as $m \to \infty$, to guarantee consistency we must have that

1. $\text{diameter}\Big(\big\{\boldsymbol{x}' \in \mathbb{R}^d \,:\, \ell(\boldsymbol{X}) = \ell(\boldsymbol{x}')\big\}\Big) \to 0$
2. $N_{\ell(\boldsymbol{X})} \to \infty$

where both conditions must hold when $m \to \infty$ and in probability with respect to the random draw of $\boldsymbol{X}$. In other words, the tree must grow unboundedly but not too fast.

In practice, a consistent algorithm may not be preferred over a nonconsistent one, see Figure 1. This due to the fact that the rate of convergence to the Bayes risk of a consistent algorithm can be arbitrarily slow, as shown by the following result.

Theorem 1 (No Free Lunch)**.** *For any sequence $a_1, a_2, \ldots$ of positive numbers converging to zero and such that $\frac{1}{16} \geq a_1 \geq a_2 \geq \cdots$ and for any consistent learning algorithm A for binary classification with zero-one loss, there exists a data distribution $\mathcal{D}$ such that $\ell_{\mathcal{D}}(f^*) = 0$ and $\mathbb{E}\big[\ell_{\mathcal{D}}\big(A(S_m)\big)\big] \geq a_m$ for all $m \geq 1$.*

Theorem 1 does not prevent a consistent algorithm from converging fast to the Bayes risk for

specific distributions $\mathcal{D}$. What the theorem shows is that if A converges to the Bayes risk for any data distribution, then it will converge arbitrarily slow for some of these distributions.

For binary classification, we can summarize the situation as follows.

- Under no assumption on η, there is no guaranteed convergence rate to Bayes risk.
- Under Lipschitz assumptions on η, the typical convergence rate to Bayes risk is $m^{-1/(d+1)}$.
- Under no assumptions on η, the typical convergence rate to the risk of the best predictor in a parametric (or finite) class $\mathcal{H}$ is $m^{-1/2}$, exponentially better than the nonparametric rate.

Note that the convergence rate $m^{-1/(d+1)}$ implies that to get ε-close to Bayes risk, we need a training set size m of order $\varepsilon^{-(d+1)}$. This exponential dependence on the number of features of the training set size is known as **curse of dimensionality** and refers to the difficulty of learning in a nonparametric setting when datapoints live in a high-dimensional space.